

Cache Simulation using Fully Associative Mapping

Problem Statement

The problem involves simulating a cache system with the following specifications:

- Main Memory size: 64K words
- Cache size: 2K words
- Block size: 16 words
- Word size: 32 bits
- Processor address: 16 bits
- Cache mapping: Fully associative
- Access pattern to simulate: Spatial access, Temporal access, and Mixed access (read and write)
- The objective is to estimate cache misses, number of cache searches, and hit/miss statistics.

The assumptions made for this simulation are as follows:

- The main memory always contains the required data, and there is never a need to access the hard disk.
- The access is performed with a fully associative cache mapping scheme.
- The cache uses the Least Recently Used (LRU) policy for cache replacement.
- The memory accesses are simulated for both read and write operations.

C++ Code

The following C++ code simulates the behavior of a fully associative cache system with spatial and temporal access patterns:

Listing 1: Cache Simulation C++ Code

```
#include <iostream>
#include <vector>
#include <climits>

class CacheBlock {
public:
    bool valid;           // Valid bit, indicates if the block
                          // contains valid data
    bool dirty;           // Dirty bit, indicates if the block
                          // has been modified
    int tag;              // Tag used for identifying the block
                          // in the cache
    int lastAccessTime;   // Time of last access for LRU
                          // replacement policy
    std::vector<int> data; // Data stored in the cache block

    // Constructor to initialize a cache block
    CacheBlock(int blockSize) {
        valid = false;    // Initially, the block is invalid
        dirty = false;    // Initially, the block is not dirty
        tag = -1;         // Initially, no tag
        lastAccessTime = 0; // Initially, no access time
        data.resize(blockSize); // Resize the data vector to hold
                               // the block's data
    }
};

class Cache {
private:
    std::vector<CacheBlock> cache; // Vector holding the cache
                                   // blocks
    int numBlocks;                 // Number of blocks in the cache
    int blockSize;                 // Size of each block (number of
                                   // words)
    int currentTime;               // Current time for LRU tracking
    int cacheMisses;               // Counter for cache misses
    int readMisses;               // Counter for read misses
    int writeMisses;              // Counter for write misses
    int cacheSearches;            // Counter for the number of
                                   // cache searches

public:
    // Constructor to initialize the cache with the number of
    // blocks and block size
    Cache(int numBlocks, int blockSize) : numBlocks(numBlocks),
        blockSize(blockSize), currentTime(0),
        cacheMisses(0), readMisses(0), writeMisses(0),
        cacheSearches(0) {
        cache.resize(numBlocks, CacheBlock(blockSize)); // Resize
        the cache to hold the specified number of blocks
    }
};
```

```

// Method to access the cache
void access(int memoryAddress, bool write) {
    currentTime++;           // Increment current time
    cacheSearches++;         // Increment the number of cache
                             searches

    int blockOffsetBits = 4; // Since the block size is 16
                             words, 4 bits are used for the block offset
    int tag = memoryAddress >> blockOffsetBits; // Extract the
                             tag from the memory address

    int blockIndex = -1;
    bool hit = false;

    // Search through the cache for a block with the same tag
    for (int i = 0; i < numBlocks; ++i) {
        if (cache[i].valid && cache[i].tag == tag) {
            blockIndex = i;
            hit = true;
            cache[i].lastAccessTime = currentTime; // Update
                                                    the access time for the block
            break;
        }
    }

    // If there is no hit, perform a cache miss
    if (!hit) {
        cacheMisses++; // Increment the cache miss counter
        if (write) {
            writeMisses++; // Increment write misses if it's a
                           write operation
        } else {
            readMisses++; // Increment read misses if it's a
                           read operation
        }
    }

    // Find the Least Recently Used (LRU) block to replace
    int lruIndex = -1;
    int minTime = currentTime;

    // LRU replacement policy
    for (int i = 0; i < numBlocks; ++i) {
        if (!cache[i].valid || cache[i].lastAccessTime <
            minTime) {
            lruIndex = i;
            minTime = cache[i].lastAccessTime;
        }
    }

    blockIndex = lruIndex;

    // Replace the cache block with the new tag
    cache[blockIndex].valid = true;
    cache[blockIndex].tag = tag;
    if (write) {

```

```

        cache[blockIndex].dirty = true; // Mark the block
        as dirty if it's a write
    }

    cache[blockIndex].lastAccessTime = currentTime; //
    Update the last access time
} else {
    if (write) {
        cache[blockIndex].dirty = true; // Mark the block
        as dirty if it's a write hit
    }
}
}

// Method to print cache statistics
void printStats() {
    std::cout << "Cache_Misses:_" << cacheMisses << std::endl;
    std::cout << "Cache_Searches:_" << cacheSearches << std::
        endl;
    std::cout << "Cache_Hit_Rate:_" << (1.0 - (double)
        cacheMisses / cacheSearches) * 100 << "%" << std::endl;
    std::cout << "Read_Misses:_" << readMisses << std::endl;
    std::cout << "Write_Misses:_" << writeMisses << std::endl;
}
};

int main() {
    // Initializing cache parameters
    int numBlocks = 128; // Cache has 128 blocks
    int blockSize = 16; // Each block contains 16 words

    // Create a cache object
    Cache cache(numBlocks, blockSize);

    // Simulating spatial access - Read and Write
    std::cout << "Simulating_Spatial_Access_-_Read:" << std::endl;
    for (int i = 0; i < 1000; i++) {
        cache.access(i, false); // Read access
    }
    cache.printStats();

    std::cout << "Simulating_Spatial_Access_-_Write:" << std::endl;
    for (int i = 0; i < 2000; i++) {
        cache.access(i, true); // Write access
    }
    cache.printStats();

    // Simulating temporal access
    std::cout << "Simulating_Temporal_Access_-_Read:" << std::endl;
    for (int i = 0; i < 1000; i++) {
        cache.access(i, false); // Read access to addresses 0 to
        999
    }
    cache.printStats();

    std::cout << "Simulating_Temporal_Access_-_Write:" << std::endl
        ;
}

```

```

    for (int i = 0; i < 4000; i++) {
        cache.access(i, true); // Write access
    }
    cache.printStats();

    return 0;
}

```

Explanation of the Code

CacheBlock Class

The `CacheBlock` class represents a single block in the cache. It contains the following attributes:

- **valid:** A boolean that indicates whether the block contains valid data.
- **dirty:** A boolean that indicates whether the block has been modified.
- **tag:** The tag that identifies the block's memory location.
- **lastAccessTime:** The timestamp of the last access to the block used for LRU replacement.
- **data:** A vector that holds the actual data stored in the block.

The constructor initializes these attributes, with **valid** set to **false**, **dirty** set to **false**, and **tag** set to -1. The **data** vector is resized based on the block size.

Cache Class

The `Cache` class represents the cache itself and performs the caching operations. It contains:

- **cache:** A vector of `CacheBlock` objects representing the cache blocks.
- **numBlocks:** The number of blocks in the cache.
- **blockSize:** The number of words per block.
- **currentTime:** The current time used to track the least recently used block.
- **cacheMisses, readMisses, writeMisses, cacheSearches:** Counters for cache misses, read/write misses, and cache searches respectively.

The **access** method performs a memory access (either read or write) and updates the cache state accordingly. If a miss occurs, it updates the cache using the LRU policy.

The **printStats** method prints the cache statistics, including the number of misses, hits, and hit rate.